

Vouch

Enterprise AI that earns the right to act.

A control layer that lets AI agents earn authority, one task at a time.

TEAM VELOCITY GATE

NIT Karnataka, Surathkal | Accenture Innovation Challenge 2026

Autonomy should be earned, not granted.

Most enterprises pick between two bad options: give the agent broad authority and absorb the risk, or review everything and lose the speed. Vouch is a third option. Routine work goes through. Expensive decisions stay with a person.

79.4%

responses released in the
measured replay

1,466

real DeepSeek replies evaluated

193

refund actions executed autonomously

0

known-wrong refund actions released

35 ms

median measured control
latency

One control plane. Three routes. Authority scoped to the job and its consequence.

Nothing goes wrong until something acts on it.

A blanket safe-or-unsafe label ignores what is at stake. A wrong sentence in an internal draft and a wrong high-value refund are not the same problem, and they should not get the same answer.

THE OPERATING TENSION

Move faster, but only as far as the evidence reaches.

1	AI platform owner	One place to plug in, whatever model is behind it.
2	Operations owner	More cases resolved without a queue, and a cap on the downside.
3	Risk and compliance	Boundaries that hold, policy per region, a record for every call.
4	Human reviewer	The reply, the evidence, the action and the reason, on one screen.

Four steps, on a loop, around every response.

01 Sense Check the claims against the systems of record, and score grounding, privacy, policy and injection.

02 Price Work out what it costs if this is wrong: the amount, and whether it can be undone.

03 Decide Weigh that cost against what this agent has already earned on this exact job.

04 Learn Wait for what actually happened, then widen or withdraw the authority.

expected loss = probability of error x action exposure

Confidence is not permission.

The gate only ever looks at three numbers. And the record it checks them against is specific: this agent, this action, this size of money, this configuration.

P(error)

How likely it is wrong

x exposure

Cost if wrong

vs budget

What it has earned

= route

Release, check, or hand over

SMALL, REVERSIBLE ACTION

Reviewing it would cost more than the mistake. So it goes through.

LARGE OR WEAKLY PROVEN ACTION

Different size, different record. It stays supervised until that record exists.

Same machinery. Different dials.

Customer support

Balanced autonomy

CHECKS

Grounding, policy, privacy

TRUST SCOPE

Per action and value band

Internal knowledge

Evidence-first answers

CHECKS

Source support, PII, injection

TRUST SCOPE

Per domain and source set

Decision support

Human accountability

CHECKS

Bias, grounding, policy, PII

TRUST SCOPE

Per class and region

FAST

Tier 2 not invoked


ADAPTIVE

GLM for borderline decisions

DEEP


GLM for meaningful side effects

It sits between the agent and what it can touch.



One path is fast and runs on every request. The other runs overnight, closes outcomes and updates what the agent has earned.

One screen. Every decision, traceable.



What happened across 1,466 real replies.

79.4%

responses released

193

refund actions executed

0

known-wrong refunds released

RESPONSE ROUTES

1,164 released

302 retained

1,466

real replies

35

borderline cases

0.018

calibration error

35 ms

median control latency

INR 17.28L

released refund value

At 100,000 decisions a month, the queue mostly goes away.

79,400

of 100,000 monthly decisions clear without a mandatory review, at the route mix we measured

VALUE EQUATION

avoided reviews x loaded review cost
+ faster resolution value
- control cost - residual error loss

Review

How many people touch one decision

Speed

How long the customer waits


Loss

What the caught mistakes would have cost

Reuse

One governance layer, every agent

Start it watching. Let it earn the rest.



When something breaks, autonomy only goes down.

RISK

Delayed or conflicting truth

CONTROL

Wait out the evidence window, and weigh sources by strength

Over-flagging

Calibration, borderline analysis, and a forced-review sample

Sensor or provider outage

Decide conservatively, cap the wait, retry once and stop

Policy or model drift

Fingerprint the configuration, and promote only through gates

Privacy and regional variation

Fixed invariants, and policy the server owns

Automation bias

The human route carries the evidence and the arithmetic

No track record can override a fixed safety boundary.

Ninety days to replace approval queues with a measured record.

A pilot needs one operations owner, one risk owner, one system of record, a defined outcome feed, and a workflow where the actions can be undone.

Evidence decides how much autonomy an agent gets.

Rushil Jain | Mayank Tiwari | Somyak Mohanta

NIT Karnataka, Surathkal | Accenture Innovation Challenge 2026