



AI reinvention made real

Accenture Innovation Challenge 2026

VOUCH · VELOCITY GATE

Team details

TEAM NAME: Velocity Gate

PROBLEM STATEMENT 1: ControlPlane.ai



Rushil Jain (Team Leader)

College: National Institute of
Technology Karnataka, Surathkal
Stream: Electronics and
Communication Engineering
Year of graduation: 2028



Somyak Mohanta

College: National Institute of
Technology Karnataka, Surathkal
Stream: Computer Science and
Engineering
Year of graduation: 2028



Mayank Tiwari

College: National Institute of
Technology Karnataka, Surathkal
Stream: Civil Engineering
Year of graduation: 2028

EXECUTIVE SUMMARY

Autonomy should be earned, not granted.

An enterprise deploying AI agents today chooses between two costly options. Grant broad authority and absorb whatever goes wrong. Or place a person in front of every action and lose the speed the automation was bought for.

Vouch is a third option. It estimates how likely each AI response is to be wrong, weighs that against what the proposed action would cost if it were wrong, and releases only what the agent's own measured record supports.

79.4%

of AI responses were approved with no human involvement

0

incorrect refunds reached a customer

Measured across every response the system released on its own

35 ms

added to each response

About a tenth of the time it takes to blink. The check is not felt

1,466

live AI responses evaluated

Generated by a production model, then scored and routed end to end

One control layer. Three outcomes: release it, verify it further, or hand it to a person.

THE PROBLEM

A confident answer is not a verified one.

Generative AI produces text that reads as authoritative whether or not it is correct. That is tolerable while the output is only text on a screen. It stops being tolerable the moment a workflow acts on it: money moves, personal data crosses a boundary, or a person makes a decision on the strength of it.

THE OPERATING TENSION

Move faster, but only as far as the evidence reaches.

FOUR PEOPLE CARRY THE CONSEQUENCES

- | | | |
|---|----------------------------|--|
| 1 | AI platform owner | One place to connect, whatever model sits behind it. |
| 2 | Operations owner | More cases resolved without a queue, and a cap on the downside. |
| 3 | Risk and compliance | Boundaries that hold, policy that varies by region, a record for every call. |
| 4 | Human reviewer | The reply, the evidence, the action and the reason, on one screen. |

WHY GENERIC CHECKING FAILS

No single checking rule fits every use case.

Each difficulty below is drawn from the brief, and each shaped a specific design decision.

THE REAL COMPLEXITY

HOW VOUCH ANSWERS IT

A chatbot, an internal copilot and a decision tool carry different risk and latency budgets

Policy is set for each use case while the underlying mechanism stays identical

Bias, hallucination and privacy overlap inside the same sentence

Signals are scored side by side and priced together, not sorted into one category

There is often no reliable real-time truth to check a claim against

Checkable claims go to the systems of record; the remainder is priced as uncertainty

Over-flagging causes alert fatigue, under-flagging creates liability

The threshold follows from the cost of being wrong rather than being chosen by hand

Agents that take actions compound risk across a conversation

Authority is granted per action and per value, never per conversation

Regulation differs by geography and continues to change

Policy is configuration, and fixed limits sit outside anything that can be earned

The model is consumed through an API, so its internals are closed

Vouch reads only inputs and outputs, so any provider can be substituted

HOW IT WORKS

Four steps, on a loop, around every response.

- | | | |
|----|---------------|--|
| 01 | Sense | Check the stated facts against the systems of record, and score grounding, privacy, policy and prompt injection. |
| 02 | Price | Work out what it costs if this is wrong: the amount at stake, and whether the action can be undone. |
| 03 | Decide | Weigh that cost against what this agent has already earned on this exact job at this size of money. |
| 04 | Learn | Wait for the confirmed outcome, then widen or withdraw the authority accordingly. |

expected cost = probability of error x cost of the action

A DECISION, IN FULL

Confidence is not permission.

The gate reads three numbers and returns one route. Nothing else influences the decision.

P(error)

How likely it is wrong

x exposure

Cost if it is wrong

vs budget

What it has earned

= route

Release, verify, or hand over

THE SAME REQUEST, AT TWO AMOUNTS

REFUND OF INR 60

The chance of being wrong is 0.2 per cent, the amount is small and the refund can be reversed. This agent has a proven record at this size, so the expected cost lands inside it.

Released automatically. No person involved.

REFUND OF INR 1,96,485

The reply is written with exactly the same confidence. But a refund this large belongs to a separate record, and the expected cost runs past what the agent has earned there.

Sent to a person, with the evidence attached.

WHAT IS NEW

Four departures from the usual safety layer.

1 Risk is priced, not labelled

A conventional checker returns safe or unsafe. Vouch multiplies the chance of error by the cost of the action, so an identical response can be released in one context and held in another.

2 Authority is earned from outcomes

Permissions are not configured by an administrator. They are a measured record, held separately for each agent, action and value band, and they shrink when results get worse.

3 The costly check is bought only when it can change the answer

A second, slower model is invoked for genuinely borderline cases rather than for every response, so verification cost and delay track difficulty instead of volume.

4 Failure always reduces autonomy

Any component that goes missing makes the system more conservative, never less. An outage cannot silently widen what the agent is allowed to do.

POLICY BY USE CASE

Same machinery. Different dials.

The brief assumes an enterprise running several AI use cases at once. Each gets its own risk tolerance and latency budget without a separate system being built for it.

Customer support

Customer facing, acts on money

CHECKS

Grounding, policy, privacy

LATENCY POSTURE

Fast: checks overlap generation

TRUST SCOPE

Per action and value band

Internal knowledge

Employee facing, mostly reading

CHECKS

Source support, PII, injection

LATENCY POSTURE

Fastest: no action to price

TRUST SCOPE

Per domain and source set

Decision support

Regulated, a person is accountable

CHECKS

Bias, grounding, policy, PII

LATENCY POSTURE

Slower by design: deeper checks


TRUST SCOPE

Per decision class and region

Verification depth is itself a setting: skip the second model, call it only for borderline cases, or apply it to every action that has a consequence.

ARCHITECTURE

It sits between the agent and what it can touch.



Inline, not after the fact

The check happens before the action reaches a system, so a wrong answer is stopped rather than reported.

Checks run beside generation

Risk scoring overlaps the model writing its reply, so most of the work costs no extra waiting time.

Every decision is recorded

The response, evidence, action, reason, provider path and latency are written to an audit log that cannot be edited in place.

MEASURED EVIDENCE

What the prototype did, and what it means.

A production AI model answered 1,466 real customer requests about refunds. Vouch scored and routed every one of them. The model can be replaced without changing anything below.

79.4%	of responses released with no human involvement	The share of work that stops consuming review capacity
0	incorrect refunds released to a customer	The autonomy above was not bought by letting errors through
INR 17,27,616	of refunds paid out under earned authority	Real money moved by the agent alone, inside its proven record
0.018	calibration error	When it says an answer is 5% likely wrong, it is wrong about 5% of the time
35 ms	added to a response, at the middle of the range	Imperceptible to a customer; the model itself takes seconds
302	responses held back for a person	Concentrated where the consequence or the uncertainty was largest

Reported alongside the two acceptance checks the prototype does not yet pass: borderline-case coverage and top-band sample size are still short.

SCALE

The design scales on volume, use cases and cost.

79,400

of 100,000 monthly decisions clear without a mandatory review, at the route mix measured here. Human attention returns to the remaining 20,600.

VALUE EQUATION

avoided reviews x loaded review cost

+ value of faster resolution

- control cost - remaining error loss

WHAT MAKES IT SCALE

Volume

Most checks are deterministic or run beside generation, so cost per decision stays roughly flat as traffic grows.

Use cases

A new workflow is a policy entry, not a new system. The gate, ledger and audit trail are shared.

Spend

The expensive second model is reserved for borderline cases, so verification cost tracks difficulty rather than volume.

Governance

One audit trail and one policy layer serve every agent, so compliance effort is paid once.

RISKS AND CONTROLS

When something breaks, autonomy only goes down.

RISK

CONTROL

Truth arrives late or conflicts

Wait out the evidence window, and weigh sources by strength

The system flags too much

Calibration, borderline analysis, and a forced-review sample

A sensor or provider fails

Decide conservatively, cap the wait, retry once and stop

Policy or the model drifts

Fingerprint the configuration, and promote only through gates

Privacy and regional rules vary

Fixed invariants, and policy the server owns

Reviewers begin rubber-stamping

The human route carries the evidence and the arithmetic

No track record can override a fixed safety boundary.

ADOPTION

Start it watching. Let it earn the rest.



WEEKS 1-2

Connect

One agent, one system of record, one outcome feed. It decides, and nothing acts on those decisions.



WEEKS 3-6

Validate

Check the probabilities against what actually happened, and put the reviewer screen in front of reviewers.



WEEKS 7-10

Constrain

Let it act on small reversible things, where the record already clears the promotion gates.



WEEKS 11-13

Expand

Open one more value band. Decide whether a second use case is worth doing.

Calibration, retained loss, review reduction, latency and operator acceptance.

PILOT GATE

Thank you

VOUCH

Our build of ControlPlane.ai. To vouch for someone is to put your name behind what they are about to do. This does it action by action, on evidence.