

Vouch

Enterprise AI that earns the right to act.

A consequence-aware control plane for measured autonomy.

TEAM VELOCITY GATE

NIT Karnataka, Surathkal | Accenture Innovation Challenge 2026

Autonomy must remain proportional to evidence.

Enterprises often choose between unsafe automation and permanent human approval. Vouch makes autonomy conditional on evidence, allowing routine workflows to proceed while high-consequence decisions remain controlled.

79.4%

responses released in the
measured replay

1,466

real DeepSeek replies evaluated

193

refund actions executed autonomously

0

known-wrong refund actions released

35 ms

median measured control
latency

One control plane. Three routes. Authority scoped to the specific workflow and consequence.

A fluent answer becomes a business risk when a workflow acts on it.

Global safe or unsafe labels ignore consequence. An incorrect internal draft and an incorrect high-value refund require different controls.

THE OPERATING TENSION

Faster execution with authority bounded by verified evidence.

1	AI platform owner	One integration surface across models and agents.
2	Operations owner	More straight-through resolution with bounded loss.
3	Risk and compliance	Fixed boundaries, regional policy and auditability.
4	Human reviewer	Response, evidence, action and reason in one record.

Four control functions convert model output into accountable action.

- 01 Sense** Verify claims and score grounding, privacy, policy, injection and secrets.
- 02 Price** Convert the proposed action into exposure using consequence and reversibility.
- 03 Decide** Compare calibrated expected loss with authority earned for this specific workflow.
- 04 Learn** Join verified outcomes to the record and expand or revoke authority.

$\text{expected loss} = \text{probability of error} \times \text{action exposure}$

Confidence never grants unlimited permission.

The gate consumes three numbers. Trust is scoped by agent, action, value band and configuration.

P(error)

Calibrated wrongness

+

x exposure

Cost if wrong

+

vs budget

Authority already earned

+

= route

Release, verify or retain

SMALL, REVERSIBLE ACTION

Eligible for initial autonomous execution when review cost exceeds total exposure.

LARGE OR WEAKLY PROVEN ACTION

Uses a different trust row and stays supervised until its own record matures.

Policy adapts to each workflow while the control mechanism remains stable.

Customer support

Balanced autonomy

CHECKS

Grounding, policy, privacy

TRUST SCOPE

Per action and value band

Internal knowledge

Evidence-first answers

CHECKS

Source support, PII, injection

TRUST SCOPE

Per domain and source set

Decision support

Human accountability

CHECKS

Bias, grounding, policy, PII

TRUST SCOPE

Per class and region

FAST

Tier 2 not invoked


ADAPTIVE

GLM for borderline decisions

DEEP

GLM for meaningful side effects

A control plane between agents and the systems they can affect.



The request path is fast and bounded. The evidence path closes outcomes, recomputes trust and monitors drift separately.

A unified interface provides decision-level traceability.



Measured results distinguish released responses from retained control.

79.4%

responses released

193

refund actions executed

0

known-wrong refunds released

RESPONSE ROUTES

1,164 released

302 retained

1,466

real replies

35

borderline cases

0.018

calibration error

35 ms

median control latency

INR 17.28L

released refund value

At enterprise scale, earned autonomy removes the review bottleneck.

79,400

monthly decisions cleared without mandatory review
at 100,000 decisions and the current route mix

VALUE EQUATION

avoided reviews x loaded review cost
+ faster resolution value
- control cost - residual error loss

Review

Touches per resolved decision

Speed

Time to customer outcome


Loss

Cost of wrong actions retained

Reuse

One governance layer across
agents

Adoption begins in shadow mode and expands by verified boundary.



Calibration, retained loss, review reduction, latency and operator acceptance.

PILOT GATE

Failure should reduce autonomy, never increase it.

RISK

Delayed or conflicting truth

CONTROL

Evidence horizons and source-weighted resolution

Over-flagging

Calibration, borderline-case analysis and forced-review sampling

Sensor or provider outage

Conservative redecision, bounded timeout and one re-entry

Policy or model drift

Configuration fingerprints and controlled promotion

Privacy and regional variation

Hard invariants and server-owned policy profiles

Automation bias

Full evidence and mathematical reason in the human route

No track record can override a fixed safety boundary.

A 90-day pilot converts static AI approvals into a measured operating model.

A pilot requires one operations owner, one risk or compliance owner, one system of record, a defined outcome feed and a workflow with reversible actions.

Verified operational evidence should determine AI autonomy.

Rushil Jain | Mayank Tiwari | Somyak Mohanta

NIT Karnataka, Surathkal | Accenture Innovation Challenge 2026